

TensorFlow Part 1

Foundation of Deep Learning

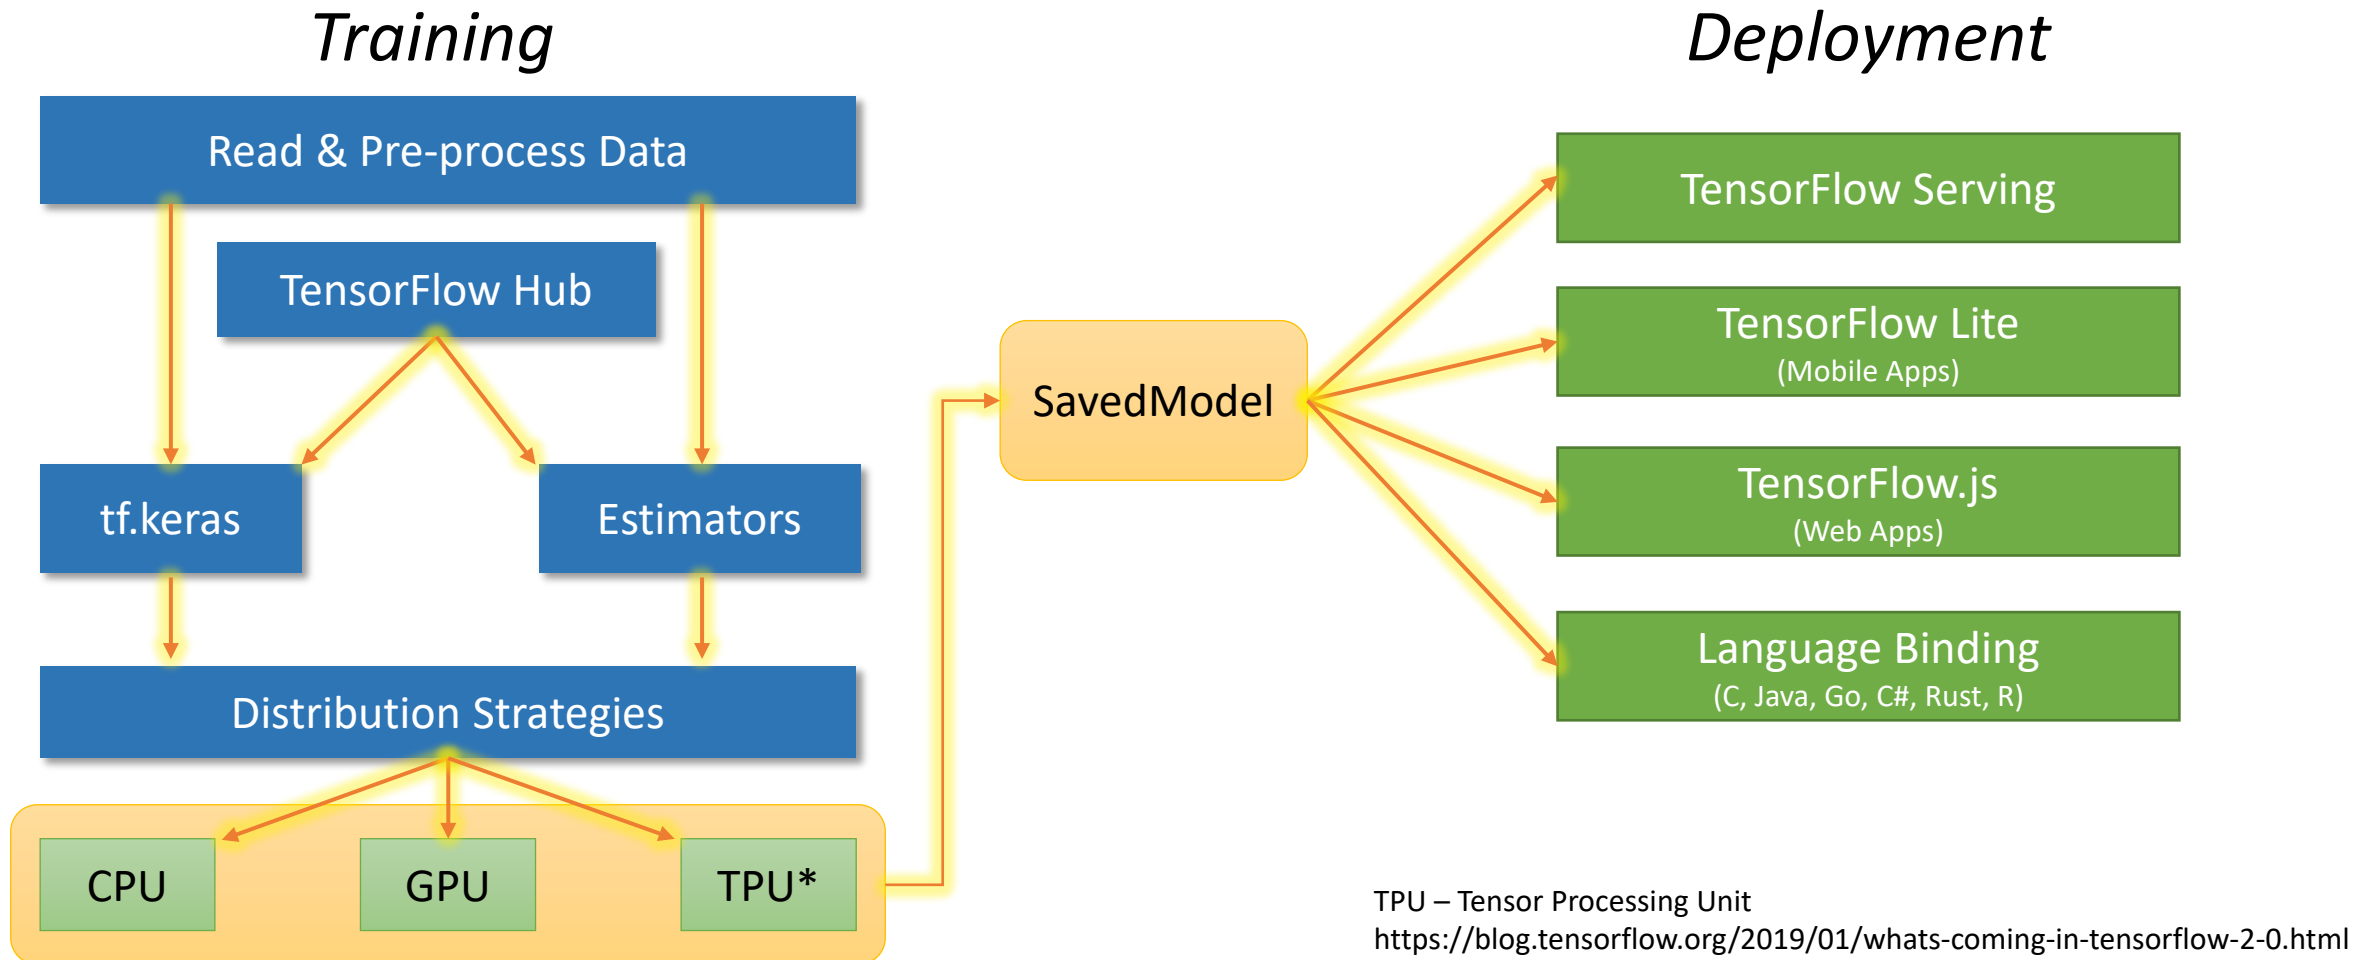
Deep Learning Frameworks

- Theano – considered grandfather of DL frameworks, developed by the Universite de Montreal in 2007
- TensorFlow - end-to-end open-source deep learning framework developed by Google and released in 2015, scalable and widely used in industry
- PyTorch - based on Torch developed by Facebook's AI research group, open-sourced on GitHub in 2017, limited visualization
- Keras - high-level neural network API written in Python can run on top of TensorFlow, CNTK and Theano.
- Others such as Apache MXNet and Microsoft CNTK

TensorFlow

- TensorFlow is free and supports the following programming languages:
 - Python
 - JavaScript
 - C/C++
 - Java
 - Go
- Python is the preferred language. Backward compatibility is not guaranteed for languages other than Python and C. [<https://www.tensorflow.org/guide/versions>]
- As can be seen from its history, it is primarily designed for **large scale distributed training and inference**.
- Although very popular for deep learning, it is also useful for other machine learning algorithms and computations.
- We can use Low-level (`tf.function`) and High-level (Keras) TensorFlow APIs
 - In this module we will cover both the low-level and high-level TensorFlow APIs.

TensorFlow 2.0 Platform



TensorFlow Hub

- Repository for publishing and using pre-trained models
- URL: <https://www.tensorflow.org/hub>

TensorFlow Hub is a repository of trained machine learning models.

TensorFlow Hub is a repository of trained machine learning models ready for fine-tuning and deployable anywhere. Reuse trained models like BERT and Faster R-CNN with just a few lines of code.



See the guide

Learn about how to use TensorFlow Hub and how it works.



See tutorials

Tutorials show you end-to-end examples using TensorFlow Hub.



See models

Find trained TF, TFLite, and TF.js models for your use case.

```
!pip install --upgrade tensorflow_hub
```

```
import tensorflow_hub as hub
```

```
model = hub.KerasLayer("https://tfhub.dev/google/nnlm-en-dim128/2")  
embeddings = model(["The rain in Spain.", "falls",  
                    "mainly", "In the plain!"])
```

```
print(embeddings.shape)  #(4,128)
```

Tensors

- TensorFlow works exclusively on **tensors**.
 - We have actually met tensors in NumPy previously.
- Tensors are data that are constructed in multiple dimensions
 - Dimension 0: scalar
 - Dimension 1: vector
 - Dimension 2 : matrix
- The number of dimensions is called the **rank**.
- We usually refer to the dimensions as **axes**.
- The number of axes is the same as the rank. For example, a rank 3 tensor (3-D Tensor) will have 3 axes.

Tensors

1

Scalar
(0-D tensor)

1
2
3

Vector
1-D Tensor

1	2
3	4
5	6

Matrix
2-D Tensor

		7	8
1	2	9	10
3	4	11	12
5	6		

3-D Tensor

...

NOTE

Do not be confused about the **dimension**.

This is a 1-D tensor but a 3-D vector.

This is a 3-D vector because it has 3 items and can represent a point in 3-D space.

Tensors Shape

- The **shape** of a tensor is a tuple of integers indicating the length along each axis.
- Example
 - The tensor on the right is of **shape (2, 3, 4)**
 - There are 3 numbers, so it is a 3D tensor (rank 3)
 - Each number indicates how many elements there are along each axis
 - Axis 0 -> 2 elements
 - Axis 1 -> 3 elements
 - Axis 2 -> 4 elements

```
m = np.array(
    [
        [1, 2, 3, 4],
        [5, 6, 7, 8],
        [9, 10, 11, 12]
    ],
    [
        [13, 14, 15, 16],
        [17, 18, 19, 20],
        [21, 22, 23, 24]
    ]
)
```

Shape of m is (2, 3, 4)




A tuple of 3 numbers,
hence it is a 3-D tensor,
there will be 3 axes

Tensors




```
[
  [
    [1, 7], [2, 8]
  ],
  [
    [3, 9], [4, 10]
  ],
  [
    [5, 11], [6, 12]
  ]
]
```

- First element of axis 0 will be [1, 7] and [2, 8]
- First element of axis 1 will be [1, 7], [3, 9], [5, 11]
- First element of axis 2 will be [1, 2], [3, 4] and [5, 6]





```
tf.gather(tensor1, 0, axis=0)
```

Tensors



Tensors



Tensors

```
m = np.array([
    [
        [1, 7], [2, 8]
    ],
    [
        [3, 9], [4, 10]
    ],
    [
        [5, 11], [6, 12]
    ]
])
t = convert_to_tensor(m)
print(tf.gather(t, 0, axis=0)) #param - tensor, index, axis
print(tf.gather(t, 0, axis=1))
print(tf.gather(t, 0, axis=2))
```

Result

```
[[1 7]
 [2 8]]
```


```
[[ 1 7]
 [ 3 9]
 [ 5 11]]
```

```
[[1 2]
 [3 4]
 [5 6]]
```

Tensors – Axis 0

- There are certain conventions in deep learning when using tensors.
- Axis 0 is usually the **samples axis**. Each element along axis 0 is a sample of data (row in the traditional feature vector)
- Sometimes is it also called the **batch axis** or **batch dimension**. This is because data are usually processed in batches. We break samples into batches along axis 0.

```
#first batch of 128 samples.  
batch = training_data[:128]
```



ASG	AGE	LOS	CLAIM
0	23	2	3619
1	20	3	4590.6
1	29	1	3850
1	27	2	11187.4
1	20	2	4368
0	25	3	5763.8
0	26	3	5108.6
1	37	2	4582.2
0	37	2	4677.4
0	30	2	3904.6

Tensors – Features and Examples

- Deep learning has been applied to many use cases and there are many different types of data.
- So far, we have been working almost exclusively with data samples and its features, see table on the right.
- In such cases, it is sufficient to use a 2-D tensor for the data
- Deep learning has been very successful especially in areas like **time series**, **images** and even **video** sequences. So a 2D tensor is not sufficient, the following slides show how data are represented using tensors.

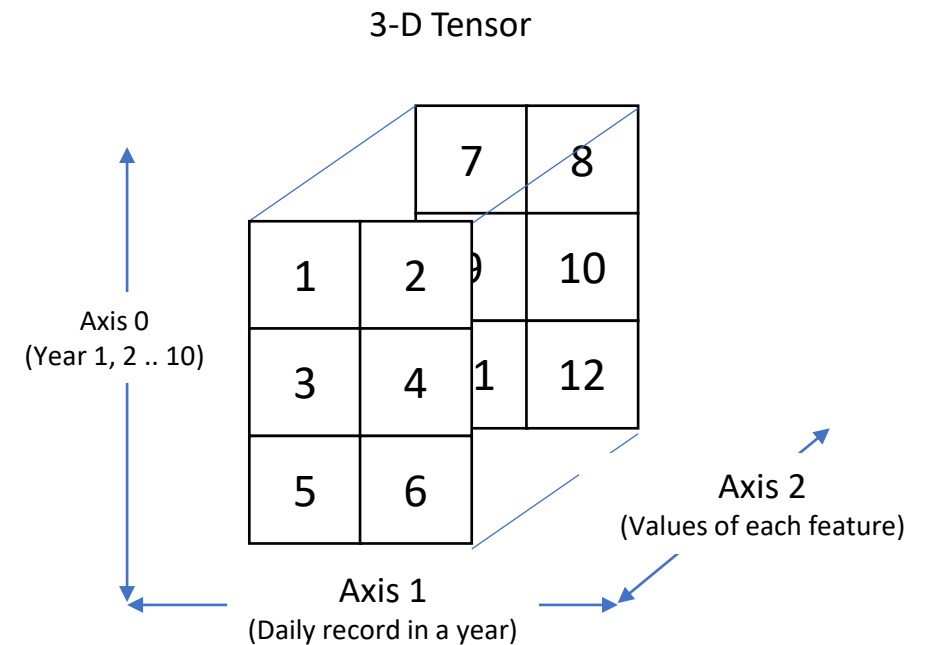
Age	Income	Years Employed
22	2500	2
45	5000	6
60	6000	20


$$\begin{bmatrix} 22 & 2500 & 2 \\ 45 & 5000 & 6 \\ 60 & 6000 & 20 \end{bmatrix}$$

2-D Tensor

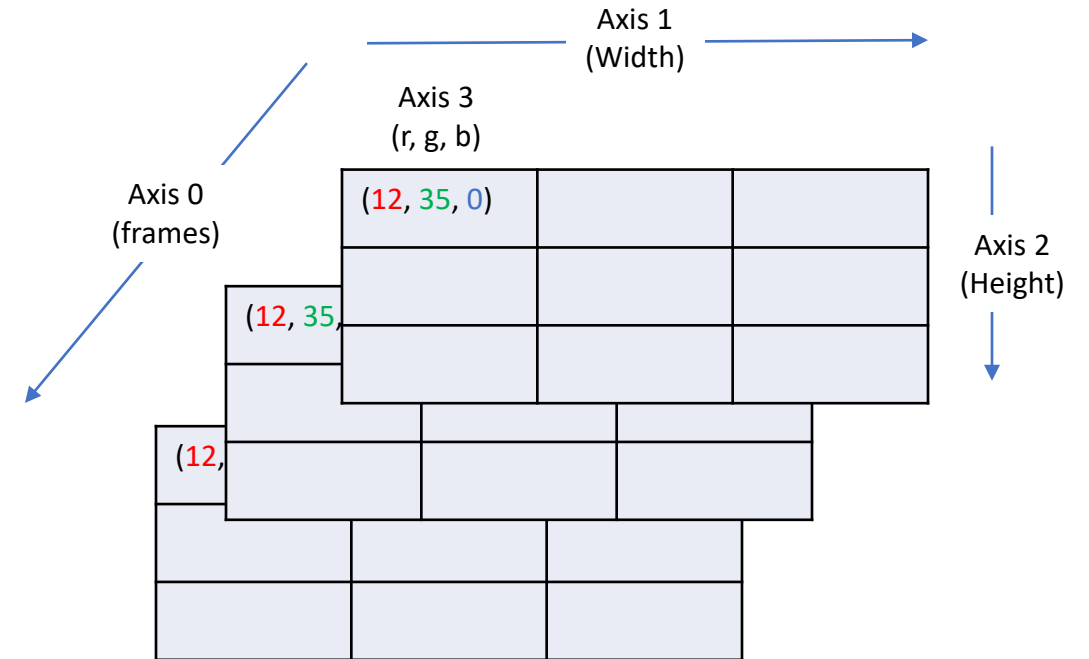
Tensors – Time Series

- Time series data uses 3D tensors
 - axis 0 – samples
 - axis 1 – timesteps
 - axis 2 - features
- Timestep refers to fixed duration, e.g. every 5 minutes.
- Example: if we store daily record of temperature (high, low, average) over 10 years, we will have a tensor of shape (10, 365, 3)
 - *Samples* - We have **10** samples where each sample is one year's worth of daily record.
 - *Timestep* - Each year we have **365** daily record.
 - *Features* - Each daily record has **3** features (high, low, average).



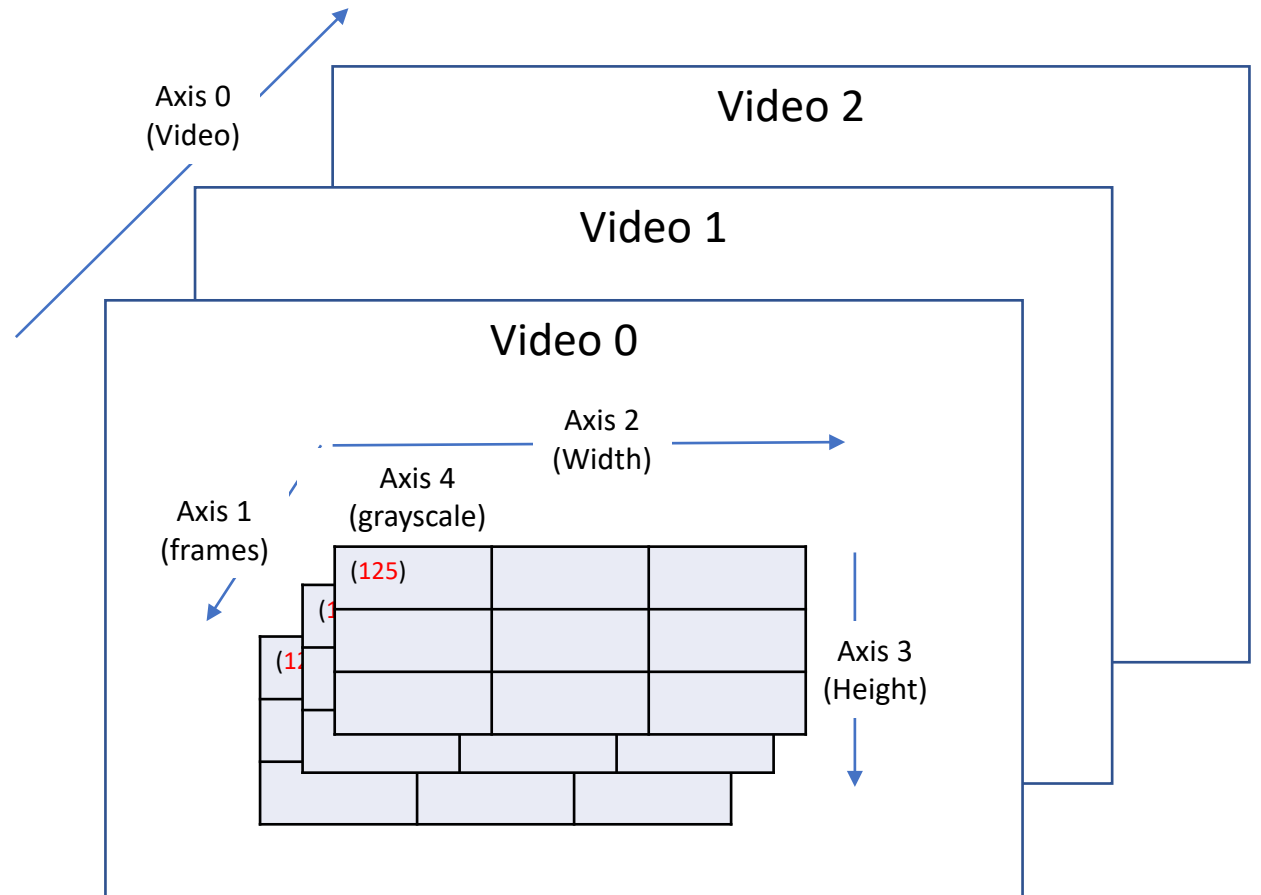
Tensors – Image Data

- TensorFlow uses 4D tensors to represent image data (multiple images).
 - Axis 0 – samples
 - Axis 1 – width,
 - Axis 2 – height
 - Axis 3 – channels. Channel refers to the colour information of each pixel.
- Note that other platforms might use different format for image data. For example, *Theano* uses a different format - samples, channels, width, height
- Each sample is a frame of image
- Each frame has width and height (number of pixels)
- Each channel can be 3 values (red, green and blue).
- Example
 - If you have 100 images, each image has resolution 1024 by 768 pixels. Each pixel has 3 values (red, green and blue).
 - Shape will be (100, 1024, 768, 3)



Tensors – Video Data

- Video can be considered a collection of image frames.
- 5D tensors (samples, frame, width, height, channels)
- Example
 - If you have 10 videos, each video is 10 seconds long and plays at 30 frames per second.
 - Each frame has resolution 176 by 144 pixels. Each pixel is greyscale with values from 0 to 255.
 - Shape will be (10, 300, 176, 144, 1)



Types of Tensors

- In TensorFlow, we can create tensors using either a **constants** or a **variable**.
- *Constants* are tensors and the values are not meant to change during the training process.
- Variables acts very much like a tensor. Strictly speaking, it is actually a type of data structure backed by a tensor. We use it to hold data that will vary during the training process.
 - If you use `tf.debugging.is_numeric_tensor()` on a variable, it will return false.
- We need to distinguish the difference between **trainable** and **non-trainable** variables.
- Trainable variables
 - The variables that the training process is trying to find the best value of.
 - For example, the β values in the linear regression algorithm ($y = \beta_1 x_1 + \beta_2 x_2 + \dots \beta_n x_n$) or the weights in a neural network.
- Non-Trainable variables
 - Other variables
 - For example, the hyper-parameter k used in k -NN.

Examples - Constants

```
import numpy as np
import tensorflow as tf
m = np.array([[[1, 7], [2, 8]], [[3, 9], [4, 10]], [[5, 11], [6, 12]]])
c = tf.constant(m)
print(c)
```

When you create a constant using TensorFlow, it is no longer a Numpy array, it will become a Tensor.

```
tf.Tensor(
[[[ 1  7]
  [ 2  8]]

[[ 3  9]
  [ 4 10]]

[[ 5 11]
  [ 6 12]]], shape=(3, 2, 2), dtype=int32)
```

Attributes like shape and data types are still applicable.

Example – Variables

```
import tensorflow as tf
k = tf.Variable(5, dtype=tf.int32, name="k")
param = tf.Variable(1.5, dtype=tf.float32, trainable=False, name="param1")
print(k)
print(param)
k.assign(10)
print(k)
```

```
<tf.Variable 'k:0' shape=() dtype=int32, numpy=5>
<tf.Variable 'param1:0' shape=() dtype=float32, numpy=1.5>
<tf.Variable 'k:0' shape=() dtype=int32, numpy=10>
```

We have created trainable and non-trainable variables. If not specified, a trainable variable will be created.

It will be better if we name the variables. This allows us to better identify the variables if necessary.

After creation, a variable's value can be changed using the `assign()` method but the shape and data type are fixed.

Operations

- In addition to variables and constants, dataflow graph contains **operations**.
- Any defined operations become nodes within dataflow graph.
- Operations represent units of **computation** while tensors represent **data**.
- Unlike some other deep learning libraries, TensorFlow comes with a wide set of operations and is flexible enough to be used for other purposes.
- Most of the mathematical operations are already provided by TensorFlow
 - Mathematics (Example add and multiply functions)
 - Array Manipulations
 - Control Flow
 - Statement Managements

Module	Available Python API
tf	tf.abs
	tf.acosh
	tf.add
	tf.add_n
	tf.angle
	tf.argmax
	tf.arg_min
	tf.asinh
	tf.assign
	tf.assign_add
	tf.assign_sub
	tf.atan
	tf.atan2
	tf.atanh
	tf.batch_to_space
	tf.batch_to_space_nd
	tf.broadcast_dynamic_shape
	tf.broadcast_static_shape

TensorFlow Operations

- TensorFlow operations has different syntax to achieve the same effect:

```
import tensorflow as tf

var1=tf.Variable([1.,2.],tf.float32, name='var_1')
var2=tf.Variable([3.,4.],tf.float32, name='var_2')

print(tf.add(var1, var2))
print(var1 + var2)
```

The result is the same:

```
tf.Tensor([4. 6.], shape=(2,), dtype=float32)
tf.Tensor([4. 6.], shape=(2,), dtype=float32)
```

TensorFlow Operations

- Very often, when constructing operations in TensorFlow, we need to work with tensor object instead of data structures like Numpy arrays, in such cases, we can use the `tf.convert_to_tensor()` function to perform the conversion.
- Example of `convert_to_tensor()`:

```
python_list = [1, 2, 3]
numpy_array = np.array([4, 5, 6])

tensor1 = tf.convert_to_tensor(python_list)
tensor2 = tf.convert_to_tensor(numpy_array)
```

TensorFlow `tf.function`

- In TensorFlow 2.x, Google encourages the use of the `tf.function` annotation when training our model.
- After we have debugged and ensured that the training functions are working properly, we should use the `tf.function` annotation for the computational intensive training process.
- The `tf.function` annotation is a tool that helps us generate Python-independent dataflow graphs out of our Python code, it uses the **AutoGraph** feature of TensorFlow.
- Basically AutoGraph converts `tf.function` annotated Python function into dataflow graph.
- These graphs are wrapped around by a *concrete function* that allows us to use the graph as if it is a function.
- The `tf.function` annotation also helps us create portable models, and it is **required** to use *SavedModel*.

TensorFlow `tf.function`


Use the `@` symbol to annotate the Python function as a `tf.function`

`@tf.function`

```
def add(num1, num2):  
    return num1 + num2
```

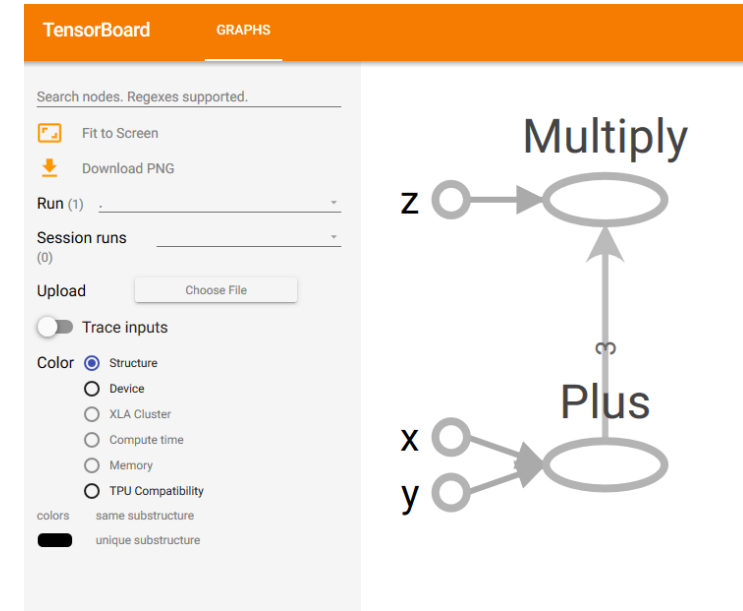
```
x1 = tf.constant([1, 2])  
x2 = tf.constant([3, 4])  
add(x1, x2)
```

The annotated Python function will be compiled into a dataflow graph



TensorBoard

- Very often, training with large amount of data can take a long time.
- It can be instructive for us to view how values change over time during the training phase.
- Google provides the TensorBoard for **visualization** and makes using TensorFlow easier.
- With TensorBoard, we can see
 - Flowchart of the layout of the dataflow graph
 - How the data flows and are processed
 - How the operations are connected
 - Summary logs
 - Performance traces

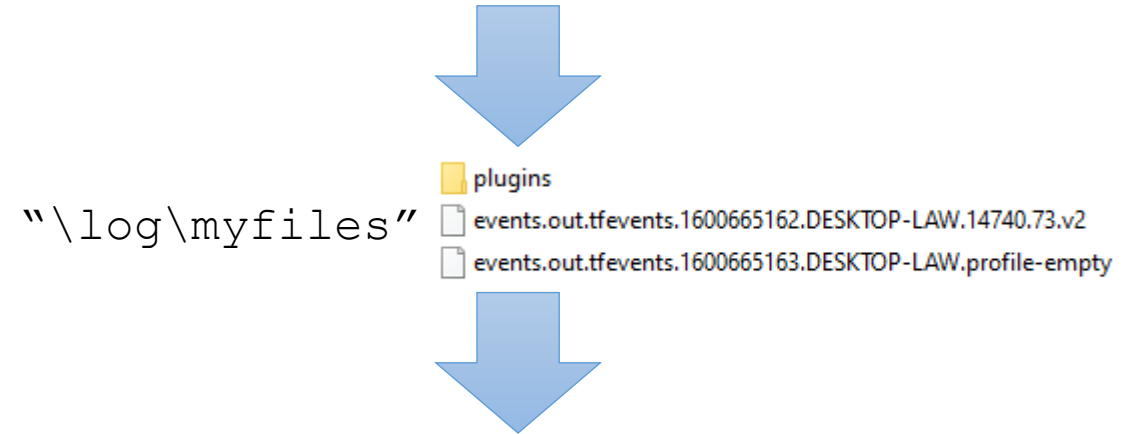


The graph created by the codes in the previous slide is visualized as shown:

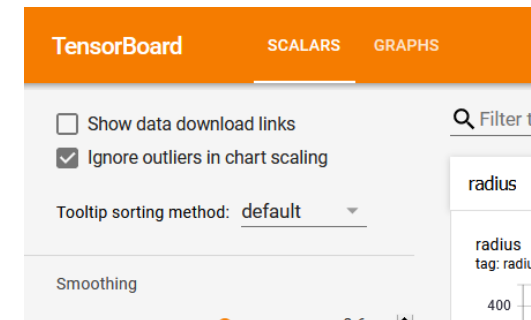
TensorBoard

- TensorBoard also allows us to visualize scalar, histogram, images and even precision-recall curves using TensorBoard.
- In our TensorFlow codes, we write data to a log directory.
 - Data to be visualized are written with the help of the `tf.summary` module.
- We then run the TensorBoard pointing to the same log directory.
- On anaconda, you need to **stop** your Jupyter Notebook kernel before running TensorBoard. Otherwise, as the Jupyter Notebook is still holding on to the file, TensorBoard will not be able to load the data!

```
log_dir = "\\log\\myfiles"
writer = tf.summary.create_file_writer(log_dir)
with writer.as_default():
    tf.summary.scalar("myvalues", 100, step)
```



```
Tensorboard --logdir \\log\\myfiles
```



Optimizers

- TensorFlow optimizers allow us to derive parameters that minimize a certain **loss** or **cost functions**.
- Some optimizers included are:
 - Gradient Descent with Momentum
 - Adadelta
 - Adagrad Dual Averaging
 - Proximal Adagrad
 - Adam
 - Adamax
 - NAdam
 - FTRL

Optimizers

To use optimizers, we

1. Define the cost function
2. Instantiate an instance of the required optimizer
3. Set the relevant parameters
4. Call the `minimize()` function iteratively.

Defines a cost function

Creates an optimizer

```
loss = lambda: (actual - predicted)**2
opt = tf.keras.optimizers.SGD(learning_rate=0.01)

for i in range(10):
    opt.minimize(loss, [weight, bias])
```

Calls the `minimize()` function

List of variables that will be updated to minimize the loss function.

Saving Models

- Sometimes, we might need to save our models either during training or for deployment.
- We can either do a *checkpoint* or *SavedModel*.
- Checkpoint
 - Captures all parameters used by a model.
 - Does not contain any description of the computation defined by the model
 - Only useful when source code that will use the saved parameter values is available.
- SavedModel
 - Includes the checkpoint data and a serialized description of the computation defined by the model
 - Does not need the source codes in order to perform inference, and thus useful for deployment.

Keras Checkpoint/SavedModel

- The `tf.keras.Model` provides
 - A `save_weights()` function that makes it easy to save a checkpoint.
 - A `save()` function that makes it easy to save as a TensorFlow SavedModel.

```
class Net(tf.keras.Model):  
    """A simple linear model."""  
  
    def __init__(self):  
        super(Net, self).__init__()  
        self.l1 = tf.keras.layers.Dense(5)  
  
    def call(self, x):  
        return self.l1(x)  
  
net = Net()  
  
net.save_weights('myfile')  
  
net.save('myfile', save_format='tf')
```

Save as
checkpoint

Save as
SavedModel

Manual Checkpoint

- We can use `tf.train.Checkpoint` to manually write variables (models data are stored as variables) to files.
- The two functions to use are `save()` and `restore()`.
- An example is shown on the right to illustrate the use of the two functions.
- The `assert_consumed()` function checks if the loading is complete and there are no errors (e.g. mismatch variables)

```
import tensorflow as tf
import os

#var is the variable to save
var = tf.Variable(0, name="var")


#Create a checkpoint
#and indicate the variable to track (i.e. var)
checkpoint = tf.train.Checkpoint(var=var)

#Save
checkpoint.save("/tmp/cp1")

#Restore
status = checkpoint.restore("/tmp/cp1").assert_consumed()
```


TensorFlow Serving

- After we have trained our model, we want to use it for prediction.
- One possibility is to use TensorFlow Serving
 - Designed for production environments
 - Makes it easy to deploy new algorithms and experiments
 - Provides out-of-the-box integration with TensorFlow models
 - Can be easily extended to serve other types of models and data.
- Refer to https://www.tensorflow.org/tfx/serving/serving_basic for tutorial on deploying and serving a TensorFlow model.



Summary

- Deep Learning Frameworks
- Tensorflow Platform and Hub
- Tensors, dimension, ranks and axes
- Trainable vs Non-trainable variables
- Tensorflow Operations
- The `tf.function` annotation
- TensorBoard
- Optimizers
- Saving and Deployment